

PAC 2017

Chao Gao

December 4, 2017


Overview

- ① About PAC
- ② About Parallel Programming
- ③ About This Travel

what is the PAC?

Parallel Application Challenge.

KNL(Knights Landing): 2nd Generation Intel Xeon Phi Processor.



Chip: 36 Tiles interconnected by 2D Mesh

Tile: 2 Cores + 2 VPU/core + 1 MB L2

Memory: MCDRAM: 16 GB on-package; High BW

DDR4: 6 channels @ 2400 up to 384GB

IO: 36 lanes PCIe Gen3. 4 lanes of DMI for chipset

Node: 1-Socket only

Fabric: Omni-Path on-package (not shown)

Vector Peak Perf: 3+TF DP and 6+TF SP Flops

Scalar Perf: ~3x over Knights Corner

Streams Triad (GB/s): MCDRAM : 400+; DDR: 90+

How fast

We want to make the program run fast, so we parallelize it.

Two big item.

- Data. How to get the data fast?
- Task. How to accomplish the task fast?

Based on specific hardware

C/C++ ([*https://github.com/memkind](https://github.com/memkind))

Allocate into DDR

```
float    *fv;  
fv = (float *)malloc(sizeof(float)*100);
```



Allocate into MCDRAM

```
float    *fv;  
fv = (float *)hbw_malloc(sizeof(float) * 100);
```

Using cache gracefully

Fetch the data fits in the cache size. Reduce the cache miss rate.

```
1 for(j = 0; j < 100; j++)  
2     for(i = 0; i < 5000; i++)  
3         x[i][j] = 2 * x[i][j];
```

```
1 for(i = 0; i < 5000; i++)  
2     for(j = 0; j < 100; j++)  
3         x[i][j] = 2 * x[i][j];
```

Vectorization

Vectorization is a special case of automatic parallelization, where a computer program is converted from a scalar implementation to a vector implementation, which processes one operation on multiple pairs of operands at once. That is SIMD (Single Instruction Multiple Data).

```
1 #pragma simd
2 for (i=0; i<n; i++)
3     c[i] = a[i] + b[i];
```


OpenMP

OpenMP is an application programming interface(API) that supports multi-platform shared memory multiprocessing programming in C, C++ and Fortran.

```
1 #pragma omp parallel for
2 for (int i = 0; i < 100000; i++) {
3     a[i] = 2 * i;
4 }
```

Data dependency causes problems

```
1 #pragma omp parallel for
2 for (int i = 2; i < 10; i++) {
3     factorial[i] = factorial[i-1];
4 }
```

```
1 #pragma omp parallel for num_threads(2)
2 for (int i = 2; i < 10; i++) {
3     factorial[i] = factorial[i-4];
4 }
```

MPI

MPI: Message Passing Interface. Data is moved from the address space of one process to that of another process through cooperative operations on each process. Mostly used in clusters of distributed memory.

When using MPI, we should partition the data manually, and have to manage how processes communicate, and have to deal with the synchronization problem.

MPI

```
1 #include "mpi.h"
2 int main(int argc, char *argv[]) {
3     int myid, numprocs;
4     MPI_Init( &argc, &argv );
5     MPI_Comm_rank( MPI_COMM_WORLD, &myid );
6     MPI_Comm_size( MPI_COMM_WORLD, &numprocs );
7     printf("I am %d of %d\n", myid, numprocs );
8     MPI_Finalize();
9 }
```

- MPI_Comm_size: number of processes
- MPI_Comm_rank: get current process ID

I have to know who am I.

```
1 int main(int argc, char *argv[]) {
2     int numprocs, myid, source; MPI_Status status;
3     char message[100];
4     MPI_Init(&argc, &argv);
5     MPI_Comm_rank(MPI_COMM_WORLD, &myid);
6     MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
7     if(myid!=0) {
8         strcpy(message, "Hello World!");
9         MPI_Send(message, strlen(message)+1, MPI_CHAR,
10                  0, 99, MPI_COMM_WORLD);
11     } else { /* myid == 0 */
12         for (source = 1; source < numprocs; source++) {
13             MPI_Recv(message, 100, MPI_CHAR, source, 99,
14                      MPI_COMM_WORLD, &status);
15             printf("%s\n", message); }
16 }
```

Algorithm

Of course, algorithm-level optimization is significant.

- Parameter tuning
- Just use a more powerful algorithm.





第五届“英特尔杯”

2017全国并行应用挑战赛颁奖晚宴

The 5th "Intel Cup" Parallel Application Challenge-PAC2017

智引天下 算出精彩

指导单位：中国计算机学会高性能计算专业委员会
主办单位：教育部高等学校计算机类专业教学指导委员会
英特尔（中国）有限公司
承办单位：北京并行科技股份有限公司

支持单位：国家超级计算广州中心
中科院计算技术研究院、中科院计算机网络信息中心
电子商务与电子支付国家工程实验室
百度云
机械工业出版社华章分社



Come on

Parallel Application Challenge(PAC)

Asia Supercomputer Community Student Supercomputer
Challenge(ASC)

The End